

Multi-Agent AI Solution for a Ride-Hailing Platform

Executive Summary

This solution proposes a production-ready multi-agent AI architecture for a ride-hailing platform similar to Uber. The system is designed to improve safety for both passengers and drivers, while also addressing operational efficiency, support workflows, fraud detection, and trust management.

Rather than relying on a single general-purpose agent, the platform is separated into specialized agents with distinct responsibilities. Each agent uses its own tools and data sources, and the results are coordinated in a production workflow. This structure makes the system more transparent, more reliable, and easier to monitor, evaluate, and improve over time.

The architecture focuses on four primary capabilities:

1. Safety and risk assessment before and during the trip
2. Real-time driver behavior monitoring using vehicle telemetry and OBD data
3. Operational escalation and trip intervention for dangerous driving
4. Customer support, incident resolution, and escalation handling

This approach aligns with the production patterns taught in the course: specialized agents, grounded tools, traceability, evaluation, and scalable workflow orchestration.

1. Business Problem

Ride-hailing platforms face a combination of operational, safety, and customer experience challenges. The business must simultaneously ensure:

- Safe pick-up and drop-off experiences
- Trustworthy driver and rider behavior
- Accurate and timely trip monitoring
- Fast responses to emergencies or suspicious activity
- Transparent support workflows for complaints, refunds, and dispute handling
- Reliable operational decisions under real-time conditions

The core problem is that safety, trip quality, and customer support are not isolated tasks. They depend on different types of signals, data, and workflows. A single generic AI agent would be too broad, too

difficult to control, and too prone to incorrect or ungrounded decisions in a high-risk environment.

A production-grade system therefore needs multiple specialized agents that contribute to a complete solution, including live vehicle telemetry. In particular, dangerous driving events such as harsh acceleration, hard braking, rapid speed changes, distracted driving, or trip anomalies are not always visible from GPS alone. OBD-based monitoring adds a more direct signal from the vehicle itself and allows faster intervention when the driver is behaving unsafely.

A production-grade system therefore needs multiple specialized agents that contribute to a complete solution.

2. Target Users

The solution is designed for four main user groups:

Passengers

- Need safe, predictable, and transparent trips
- Require fast assistance in emergency or service disruption situations
- Expect clear communication during delays, cancellations, and complaints

Drivers

- Need route support, safety guidance, and protection against false complaints
- Want faster support when a trip becomes risky or ambiguous
- Need help managing operational issues without manual intervention

Customer Support and Operations Teams

- Monitor trip anomalies and risk indicators
- Review escalations and safety cases
- Manage refunds, case documentation, and policy compliance

Platform Managers

- Monitor business KPI trends such as safety incidents, driver reliability, support burden, and customer trust
- Make strategic decisions based on patterns across rides, regions, and driver groups

3. Proposed Multi-Agent System

The solution includes four specialized agents and supporting tools and data sources.

Agent 1: Safety Risk and Trip Assessment Agent

This agent evaluates the safety profile of a ride before it begins and during a trip when conditions change.

Responsibilities:

- Assess driver profile and trip history
- Evaluate rider profile and safety context
- Analyze route risk, pickup area risk, and trip patterns
- Detect suspicious or abnormal trip behavior
- Trigger escalation for high-risk scenarios

Tools and data sources:

- Driver historical data and performance records
- Rider profile and trip history
- Route risk analysis and geofencing data
- Safety policy documents and operating procedures

Example use cases:

- Driver fatigue or repeated late-night long shifts
- Area with historically high incident rates
- Mismatch between route pattern and expected trip behavior
- Repeated complaints or suspicious behavior patterns

Agent 2: Vehicle Telemetry and Dangerous Driving Agent

This agent is explicitly designed for vehicle telemetry captured via OBD-II or telematics devices. It monitors the vehicle's live state and detects dangerous driving patterns in near real time.

Responsibilities:

- Collect and evaluate OBD signals such as speed, RPM, coolant temperature, acceleration, braking force, engine load, throttle position, and steering events
- Detect harsh braking, rapid acceleration, aggressive cornering, speeding, and excessive idling
- Detect anomalies such as sudden braking, acceleration spikes, strong steering inputs, high engine load, or unsafe driving patterns
- Trigger warnings or escalate to operations when a driver enters a dangerous behavior pattern
- Support coaching and safety scoring after the trip ends

Tools and data sources:

- OBD-II vehicle telemetry stream

- Telematics integration layer from the driver app or vehicle gateway
- GPS and motion data
- Historical driver telematics scorecards
- Safety policy rules for aggressive driving thresholds

Example use cases:

- Driver brakes sharply in a congested urban area and generates multiple hard-brake events
- Vehicle exceeds speed threshold during a night trip in a restricted zone
- Excessive steering corrections or aggressive lane-change events appear in the telemetry
- Short sequence of acceleration spikes suggests dangerous or distracted driving behavior

This agent is critical because dangerous driving is often not visible from trip metadata alone. By integrating OBD and telematics data, the system can react earlier and more accurately to risky behavior that may otherwise remain hidden until a complaint or incident occurs.

Agent 3: Driver Operations and In-Trip Monitoring Agent

This agent watches live trip conditions and identifies operational or safety anomalies while the ride is underway.

Responsibilities:

- Monitor trip trajectory and route adherence
- Detect suspicious detours, long idle periods, or route deviation
- Observe driver response patterns and communication quality
- Trigger emergency workflows if needed
- Support driver safety and operational continuity

Tools and data sources:

- GPS, route, and telematics data
- Driver availability, rest-time, and shift data
- Live trip telemetry and event stream
- Safety escalation procedures and standard response playbooks

Example use cases:

- Driver leaves the route for an extended period
- Passenger reports unusual behavior or pressure
- Driver fails to respond to trip check-ins
- Trip stalls in a high-risk area or against known safety policy

Agent 4: Customer Support and Incident Resolution Agent

This agent handles customer-facing issues after the trip or during service disruption events.

Responsibilities:

- Investigate claims, delays, safety complaints, or fare issues
- Summarize trip context for human agents
- Validate policies and compensation eligibility
- Create support tickets and generate decisions or recommendations

Tools and data sources:

- Trip history and metadata
- Policy libraries, refund rules, and support playbooks
- Incident and support ticket systems
- Prior case histories and dispute decisions

Example use cases:

- Passenger reports unsafe or inappropriate driver behavior
- Trip ended unexpectedly or in the wrong location
- Driver disputes a complaint or challenge
- Delay or cancellation requires compensation

4. Why a Multi-Agent Design Is Better Than a Single Agent

A single large agent would be too general and too difficult to optimize for high-risk operational decisions. In ride-hailing systems, different tasks require different priorities, tools, and reasoning paths.


The multi-agent design is better because:

- Each agent has a clear scope and objective
- Safety decisions are separated from customer support decisions
- Real-time operational monitoring is independent from claims handling
- Vehicle telemetry adds a critical safety layer that GPS alone cannot detect
- Failures are easier to isolate and debug
- The system can apply more precise policies and escalation rules
- It is easier to add tools, evaluation criteria, and operational safeguards

This separation improves not just performance, but also trust and maintainability. In a production environment, explainability and controlled behavior are as important as accuracy.

5. Information Flow Between Agents

The agents work together as a coordinated workflow:



This flow ensures that safety is checked before a trip is accepted, then monitored during the trip using both route data and vehicle telemetry, and finally resolved with structured support logic when incidents occur.

Example workflow scenario: suspicious trip escalation

A concrete example of the operational workflow is shown below:

1. A passenger requests a ride at 11:45 PM in an area with a higher-than-normal incident rate.
2. The Safety Risk and Trip Assessment Agent checks the driver profile, rider history, route risk, and local safety conditions.

3. The system detects that the driver has been active for several hours and the route passes through a high-risk district. The risk score is elevated.
4. The Vehicle Telemetry and Dangerous Driving Agent receives OBD data and detects harsh braking, rapid acceleration, and speeding above the threshold.
5. The agent does not automatically approve the trip. Instead, it sends a warning to the operations team and asks for manual review or a safer cautionary route.
6. The trip is accepted only after the route is verified and a safety check is completed.
7. During the trip, the Driver Operations and In-Trip Monitoring Agent detects a detour and a delayed response from the driver.
8. The vehicle telemetry stream also shows a hard-brake sequence and elevated RPM near a high-risk junction.
9. The system triggers a safety alert, shares the event with the operations dashboard, and identifies the trip as requiring escalation.
10. The passenger receives a check-in message from the platform and is offered a safe fallback option if needed.
11. If the passenger later raises a complaint, the Customer Support and Incident Resolution Agent reviews trip telemetry, policy rules, and incident history to recommend a fair resolution.
12. Final output is a structured report for operations: trip risk status, dangerous driving events, support outcome, and any corrective action required for future rides.

This example shows how the workflow links risk detection, real-time safety monitoring, OBD telemetry, and customer support into one operational loop.

6. Production Readiness Plan

6.1 Observability Strategy

For this system to be trusted in production, the organization must capture operational and behavioral trace data for every agent run.

Recommended monitoring signals:

- Full trace of each agent request, including prompts, inputs, tool calls, and outputs
- Latency for each stage of the workflow
- Token usage and cost per agent run
- Tool execution results and failure rates
- Risk classification history and escalation events
- Trip anomaly logs: route deviation, delayed response, idle trip, cancellation patterns
- Vehicle telemetry signals: speed, RPM, acceleration, braking, steering inputs, engine load, hard-brake events

- Dangerous driving scorecards by driver and trip
- Support case outcomes and accountability records

This information helps the team:

- Identify which tool failed or returned unreliable data
- Understand why an agent made a certain risk decision
- Compare model versions and prompt changes against baseline performance
- Detect drift in operational patterns and route behavior
- Reduce the chance of silent failures or bad escalations
- Link dangerous driving behavior to a specific trip, route, and driver event stream

6.2 Evaluation Strategy

The solution should not be judged only by whether it responds politely. It must be measured against business and safety criteria.

Evaluation dataset should include examples such as:

- Safe trip with normal conditions
- High-risk route or suspicious pickup pattern
- Driver requests emergency support during trip
- Passenger complaint without valid cause
- Driver false accusation or dispute case
- Cancelled trip with compensation eligibility
- Fraudulent route pattern or manipulated location behavior
- Hard-brake / speeding / rapid-acceleration events captured from OBD telemetry
- Aggressive driving patterns that should trigger a coaching or safety alert

Key evaluation metrics:

- Correctness of risk assessment
- Safety alignment with policy and operational procedures
- Response quality and coherence
- Tool usage accuracy
- Groundedness against trip data and policy documents
- Safety compliance and escalation appropriateness
- Dangerous-driving detection accuracy from OBD telemetry

These evaluations can be integrated into the software lifecycle:

- Before deployment

- After prompt or model changes
- Periodically during operations
- As part of release gates for production updates

6.3 Reliability and Safety Controls

The system should include several controls to ensure consistency and safe operations:

- Versioned prompts and model configurations
- Controlled tool access with clear permissions
- Risk thresholds for escalation to human operators
- Policy-grounded reasoning using company knowledge bases
- Data filtering and privacy-safe handling of personal information
- Human review for highly sensitive decisions such as emergency cases or critical rider complaints
- Driver safety scoring based on OBD events
- Audit logs for every decisive action

This allows the AI system to be trustworthy, explainable, and maintainable over time.

7. Workflow Design

The end-to-end workflow is as follows:

Step 1: Trip Intake and Pre-Check

The platform receives a trip request or ride assignment. The Safety Risk and Trip Assessment Agent checks:

- driver qualifications and recent behavior
- rider risk profile and trip context
- route risk indicators and pickup safety conditions

Step 2: Vehicle Telemetry Monitoring

As soon as the trip begins, the Vehicle Telemetry and Dangerous Driving Agent starts evaluating OBD and telematics data:

- speed and engine RPM
- harsh braking and rapid acceleration events
- aggressive steering inputs
- excessive idling or abnormal engine load
- sudden unsafe behavior patterns

Step 3: Route and Safety Monitoring

If the trip is approved, the Driver Operations and In-Trip Monitoring Agent begins monitoring:

- route adherence
- trip duration and detours
- unusual idle time
- driver availability and communication issues

Step 4: Escalation and Safety Workflow

If the system detects a risk condition, it triggers alerts and may escalate to a human operator or support team. This protects both the passenger and the driver.

Step 5: Customer Incident Handling

If the rider or driver raises a complaint, the Customer Support and Incident Resolution Agent reviews trip evidence, policy rules, and initiative history to generate a resolution recommendation.

Step 6: Final Business Output

The platform produces a structured final result such as:

- trip accepted and monitored normally
- trip escalated for human review
- dangerous driving event recorded and coached
- complaint routed to support and investigated
- compensation or policy-based resolution recommended
- post-incident safety insights shared with operations

7. Implementation and Deployment Model

This workflow can be deployed using Microsoft Foundry as a production-ready multi-agent architecture:

- Agents are created as persistent resources
- Each agent has a specialized role and tool set
- The orchestration layer coordinates the sequence of actions
- Monitoring is connected to Application Insights or a similar telemetry system
- OBD and telematics data are connected to the in-trip monitoring pipeline
- Evaluation datasets validate quality before release
- The workflow can be exposed through an API or internal business workflow

This approach turns a prototype into a deployable operational system with measurable reliability and maintainability.

8. Why This Solution Fits the Course Concepts

This design follows the core ideas taught during the program:

- Multi-agent specialization instead of a single monolithic agent
- Use of tools and grounded data sources to reduce hallucination
- Knowledge bases for policies, safety rules, and historical incidents
- Vehicle telemetry and OBD data as situational inputs for live safety decisions
- System observability for debugging and performance monitoring
- Evaluation datasets and quality metrics to measure confidence and consistency
- Deployment-oriented orchestration that supports real business processes

This is not only an AI demo. It is a structured business solution designed to operate in a real-world ride-hailing environment with strong emphasis on trust, safety, and operational control.

9. Conclusion

The proposed solution addresses a realistic and important challenge for ride-hailing platforms: how to deliver safe, reliable, and efficient mobility services while managing operational risk and customer trust.

By using a multi-agent architecture focused on safety, live operations, vehicle telemetry, and customer support, the system provides better decision quality than a single generic agent. It is also designed for real production use through observability, evaluation, policy grounding, and operational orchestration.

In particular, the addition of OBD telemetry makes the system more capable of detecting dangerous driving in real time—an essential requirement for a modern ride-hailing platform that wants to protect both riders and drivers while maintaining a strong safety culture.

In short, this architecture shows how AI can move from concept prototype to a production-ready platform that supports both business performance and user safety.